

In [1]: `import warnings`
`warnings.filterwarnings('ignore')`

In [2]: `import tensorflow`
`from tensorflow.keras.preprocessing.image import ImageDataGenerator`

In [3]: `train_Datagen = ImageDataGenerator(rescale=1./255,`
 `shear_range=0.2,`
 `zoom_range=0.2,`
 `horizontal_flip=True)`

In [4]: `training_set = train_Datagen.flow_from_directory('train',`
 `target_size=(128,128),`
 `batch_size=32,`
 `class_mode='binary')`

Found 40 images belonging to 2 classes.

In [5]: `test_Datagen = ImageDataGenerator(rescale=1./255)`

In [6]: `testing_set = test_Datagen.flow_from_directory('test',`
 `target_size=(128,128),`
 `batch_size=32,`
 `class_mode='binary')`

Found 20 images belonging to 2 classes.

In [7]: `from tensorflow.keras.models import Sequential`
`from tensorflow.keras.layers import Dense, Conv2D, MaxPool2D, Flatten, Dropout`

In [8]: `model = Sequential()`

In [9]: `model.add(Conv2D(filters = 64,`
 `kernel_size = (5,5),`
 `activation = 'relu',`
 `padding = 'same',`
 `input_shape = [128 , 128, 3]))`

In [10]: `model.add(MaxPool2D(pool_size = (2,2), strides = 2))`

In [11]: `model.add(Conv2D(filters = 32,`
 `kernel_size = (5,5),`
 `activation = 'relu'))`

In [12]: `model.add(MaxPool2D(pool_size = (2,2), strides = 2))`

In [13]: `model.add(Flatten())`

In [14]: `model.add(Dense(units = 128,`
 `activation = 'relu'))`
`model.add(Dropout(0.4))`

In [15]: `model.add(Dense(units = 64, activation = 'relu'))`
`model.add(Dense(units = 1, activation = 'sigmoid'))`

In [16]: `model.compile(optimizer=tensorflow.keras.optimizers.Adam(),`
 `loss=tensorflow.keras.losses.BinaryCrossentropy(),`
 `metrics=['accuracy'])`

In [17]: `model.summary()`

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 128, 128, 64)	4864
max_pooling2d (MaxPooling2D)	(None, 64, 64, 64)	0
conv2d_1 (Conv2D)	(None, 60, 60, 32)	51232
max_pooling2d_1 (MaxPooling2D)	(None, 30, 30, 32)	0
flatten (Flatten)	(None, 28800)	0
dense (Dense)	(None, 128)	3686528
dropout (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 64)	8256
dense_2 (Dense)	(None, 1)	65

=====
Total params: 3,750,945
Trainable params: 3,750,945
Non-trainable params: 0
=====

In [18]: `model.fit(training_set, validation_data=testing_set, verbose=1, epochs=50)`

Epoch 1/50
2/2 [=====] - 3s 762ms/step - loss: 0.7012 - accuracy: 0.6000
- val_loss: 1.3307 - val_accuracy: 0.5000
Epoch 2/50
2/2 [=====] - 1s 397ms/step - loss: 1.3980 - accuracy: 0.4750
- val_loss: 0.7116 - val_accuracy: 0.5000
Epoch 3/50
2/2 [=====] - 1s 1s/step - loss: 0.7128 - accuracy: 0.4750 -
val_loss: 0.6949 - val_accuracy: 0.4500
Epoch 4/50
2/2 [=====] - 1s 946ms/step - loss: 0.6893 - accuracy: 0.6500
- val_loss: 0.6943 - val_accuracy: 0.4500
Epoch 5/50
2/2 [=====] - 1s 895ms/step - loss: 0.6887 - accuracy: 0.5250
- val_loss: 0.7030 - val_accuracy: 0.4000
Epoch 6/50
2/2 [=====] - 1s 350ms/step - loss: 0.6891 - accuracy: 0.5000
- val_loss: 0.7164 - val_accuracy: 0.3000
Epoch 7/50
2/2 [=====] - 1s 372ms/step - loss: 0.6735 - accuracy: 0.6000
- val_loss: 0.7339 - val_accuracy: 0.3000
Epoch 8/50
2/2 [=====] - 1s 909ms/step - loss: 0.6909 - accuracy: 0.4750
- val_loss: 0.7648 - val_accuracy: 0.5000
Epoch 9/50
2/2 [=====] - 1s 886ms/step - loss: 0.7245 - accuracy: 0.4000
- val_loss: 0.7235 - val_accuracy: 0.4500
Epoch 10/50
2/2 [=====] - 1s 361ms/step - loss: 0.6953 - accuracy: 0.5750
- val_loss: 0.7143 - val_accuracy: 0.3500
Epoch 11/50
2/2 [=====] - 1s 376ms/step - loss: 0.6670 - accuracy: 0.5500
- val_loss: 0.7162 - val_accuracy: 0.5000
Epoch 12/50
2/2 [=====] - 1s 886ms/step - loss: 0.6814 - accuracy: 0.5500
- val_loss: 0.7093 - val_accuracy: 0.5000
Epoch 13/50
2/2 [=====] - 1s 367ms/step - loss: 0.7011 - accuracy: 0.5000
- val_loss: 0.6975 - val_accuracy: 0.3500
Epoch 14/50
2/2 [=====] - 1s 362ms/step - loss: 0.6770 - accuracy: 0.6250
- val_loss: 0.6957 - val_accuracy: 0.5000
Epoch 15/50
2/2 [=====] - 1s 897ms/step - loss: 0.7205 - accuracy: 0.5500
- val_loss: 0.6977 - val_accuracy: 0.5000
Epoch 16/50
2/2 [=====] - 1s 337ms/step - loss: 0.7193 - accuracy: 0.4000
- val_loss: 0.6933 - val_accuracy: 0.6500
Epoch 17/50
2/2 [=====] - 1s 384ms/step - loss: 0.6822 - accuracy: 0.5500
- val_loss: 0.6929 - val_accuracy: 0.5000
Epoch 18/50
2/2 [=====] - 1s 897ms/step - loss: 0.6899 - accuracy: 0.5000
- val_loss: 0.6968 - val_accuracy: 0.5000
Epoch 19/50
2/2 [=====] - 1s 367ms/step - loss: 0.6889 - accuracy: 0.5000
- val_loss: 0.7005 - val_accuracy: 0.5000
Epoch 20/50
2/2 [=====] - 1s 902ms/step - loss: 0.6726 - accuracy: 0.6000
- val_loss: 0.7050 - val_accuracy: 0.4000
Epoch 21/50
2/2 [=====] - 1s 364ms/step - loss: 0.6752 - accuracy: 0.6000
- val_loss: 0.7118 - val_accuracy: 0.4500
Epoch 22/50
2/2 [=====] - 1s 382ms/step - loss: 0.6830 - accuracy: 0.5500
- val_loss: 0.7088 - val_accuracy: 0.4000
Epoch 23/50
2/2 [=====] - 1s 371ms/step - loss: 0.6867 - accuracy: 0.6000
- val_loss: 0.6971 - val_accuracy: 0.4500
Epoch 24/50
2/2 [=====] - 1s 377ms/step - loss: 0.6788 - accuracy: 0.7750
- val_loss: 0.7038 - val_accuracy: 0.5000
Epoch 25/50
2/2 [=====] - 1s 359ms/step - loss: 0.6746 - accuracy: 0.6250
- val_loss: 0.7157 - val_accuracy: 0.5000
Epoch 26/50
2/2 [=====] - 1s 882ms/step - loss: 0.7268 - accuracy: 0.4250
- val_loss: 0.7030 - val_accuracy: 0.3500
Epoch 27/50
2/2 [=====] - 1s 880ms/step - loss: 0.6678 - accuracy: 0.6250
- val_loss: 0.7023 - val_accuracy: 0.5000
Epoch 28/50
2/2 [=====] - 1s 869ms/step - loss: 0.7038 - accuracy: 0.5000
- val_loss: 0.6966 - val_accuracy: 0.5500
Epoch 29/50
2/2 [=====] - 1s 915ms/step - loss: 0.6646 - accuracy: 0.6500
- val_loss: 0.7006 - val_accuracy: 0.5000
Epoch 30/50
2/2 [=====] - 1s 895ms/step - loss: 0.6599 - accuracy: 0.6500
- val_loss: 0.7050 - val_accuracy: 0.4500
Epoch 31/50
2/2 [=====] - 1s 905ms/step - loss: 0.6588 - accuracy: 0.6500
- val_loss: 0.6980 - val_accuracy: 0.4000
Epoch 32/50
2/2 [=====] - 1s 351ms/step - loss: 0.6290 - accuracy: 0.7000
- val_loss: 0.6970 - val_accuracy: 0.5000
Epoch 33/50
2/2 [=====] - 1s 873ms/step - loss: 0.5953 - accuracy: 0.7000
- val_loss: 0.7084 - val_accuracy: 0.3500
Epoch 34/50
2/2 [=====] - 1s 348ms/step - loss: 0.6657 - accuracy: 0.6500
- val_loss: 0.7068 - val_accuracy: 0.5500
Epoch 35/50
2/2 [=====] - 1s 343ms/step - loss: 0.6615 - accuracy: 0.5250
- val_loss: 0.6927 - val_accuracy: 0.6000
Epoch 36/50
2/2 [=====] - 1s 338ms/step - loss: 0.6227 - accuracy: 0.7250
- val_loss: 0.6837 - val_accuracy: 0.5000
Epoch 37/50
2/2 [=====] - 1s 872ms/step - loss: 0.6134 - accuracy: 0.8000
- val_loss: 0.7003 - val_accuracy: 0.5500
Epoch 38/50
2/2 [=====] - 1s 863ms/step - loss: 0.6214 - accuracy: 0.6250
- val_loss: 0.7706 - val_accuracy: 0.6000
Epoch 39/50
2/2 [=====] - 1s 873ms/step - loss: 0.6098 - accuracy: 0.7000
- val_loss: 0.6789 - val_accuracy: 0.6000
Epoch 40/50
2/2 [=====] - 1s 376ms/step - loss: 0.5523 - accuracy: 0.8250
- val_loss: 0.6816 - val_accuracy: 0.5500
Epoch 41/50
2/2 [=====] - 1s 345ms/step - loss: 0.5477 - accuracy: 0.8500
- val_loss: 0.7141 - val_accuracy: 0.5500
Epoch 42/50
2/2 [=====] - 1s 922ms/step - loss: 0.5387 - accuracy: 0.7250
- val_loss: 0.7378 - val_accuracy: 0.4500
Epoch 43/50
2/2 [=====] - 1s 366ms/step - loss: 0.4727 - accuracy: 0.7750
- val_loss: 0.7521 - val_accuracy: 0.4500
Epoch 44/50
2/2 [=====] - 1s 906ms/step - loss: 0.4184 - accuracy: 0.7750
- val_loss: 0.7644 - val_accuracy: 0.5000
Epoch 45/50
2/2 [=====] - 1s 892ms/step - loss: 0.5171 - accuracy: 0.7500
- val_loss: 0.7293 - val_accuracy: 0.6500
Epoch 46/50
2/2 [=====] - 1s 369ms/step - loss: 0.4783 - accuracy: 0.8500
- val_loss: 0.7691 - val_accuracy: 0.6500
Epoch 47/50
2/2 [=====] - 1s 342ms/step - loss: 0.3470 - accuracy: 0.8750
- val_loss: 1.0359 - val_accuracy: 0.6000
Epoch 48/50
2/2 [=====] - 1s 356ms/step - loss: 0.3992 - accuracy: 0.8250
- val_loss: 1.0483 - val_accuracy: 0.5000
Epoch 49/50
2/2 [=====] - 1s 864ms/step - loss: 0.4094 - accuracy: 0.8500
- val_loss: 0.9379 - val_accuracy: 0.6000
Epoch 50/50
2/2 [=====] - 1s 922ms/step - loss: 0.3530 - accuracy: 0.8750
- val_loss: 0.7980 - val_accuracy: 0.6500

Out[18]: <keras.callbacks.History at 0x1e55d692dc0>


In [19]: `import numpy as np`

In [20]: `from tensorflow.keras.preprocessing import image`

In [21]: `test_img = image.load_img('cat.jpg', target_size=(128, 128, 3))`

In [22]: `import matplotlib.pyplot as plt`
`plt.imshow(test_img)`

Out[22]: <matplotlib.image.AxesImage at 0x1e55f0920a0>



In [23]: `test_img = image.img_to_array(test_img)`
`test_img = np.expand_dims(test_img, axis=0)`
`test_img = test_img/255`

In [24]: `result = model.predict(test_img)`

In [25]: `training_set.class_indices`

Out[25]: {'cats': 0, 'dogs': 1}

In [26]: `result`

Out[26]: array([[0.18541053]], dtype=float32)

In [27]: `if np.round(result[0][0]) == 1 :`
 `prediction = "It's Dog image.."`
`else :`
 `prediction = "It's Cat image.."`

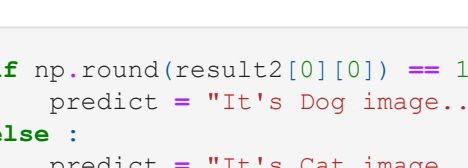
`print(prediction)`

It's Cat image..

In [28]: `test_img2 = image.load_img('dog.jpg', target_size=(128, 128, 3))`

In [29]: `plt.imshow(test_img2)`

Out[29]: <matplotlib.image.AxesImage at 0x1e55f1176a0>



In [30]: `test_img2 = image.img_to_array(test_img2)`
`test_img2 = np.expand_dims(test_img2, axis=0)`
`test_img2 = test_img2/255`

In [31]: `result2 = model.predict(test_img2)`

In [32]: `result2`

Out[32]: array([[0.8717852]], dtype=float32)

In [33]: `if np.round(result2[0][0]) == 1 :`
 `predict = "It's Dog image.."`
`else :`
 `predict = "It's Cat image.."`

`print(predict)`

It's Dog image..

In []: